

RETAIL (Capstone Project 1)

DESCRIPTION

It is a critical requirement for business to understand the value derived from a customer. RFM is a method used for analyzing customer value. Customer segmentation is the practice of segregating the customer base into groups of individuals based on some common characteristics such as age, gender, interests, and spending habits. Perform customer segmentation using RFM analysis. The resulting segments can be ordered from most valuable (highest recency, frequency, and value) to least valuable (lowest recency, frequency, and value). Dataset Description

This is a transnational data set which contains all the transactions that occurred between 01/12/2010 and 09/12/2011 for a UK-based and registered non-store online retail. The company mainly sells unique and all-occasion gifts.

Variables Description InvoiceNo Invoice number. Nominal, a six digit integral number uniquely assigned to each transaction. If this code starts with letter 'c', it indicates a cancellation StockCode Product (item) code. Nominal, a five digit integral number uniquely assigned to each distinct product Description Product (item) name. Nominal Quantity The quantities of each product (item) per transaction. Numeric InvoiceDate Invoice Date and time. Numeric, the day and time when each transaction was generated UnitPrice Unit price. Numeric, product price per unit in sterling CustomerID Customer number. Nominal, a six digit integral number uniquely assigned to each customer Country Country name. Nominal, the name of the country where each customer resides

Project Task: Week 1

1. Data Cleaning:

- Perform a preliminary data inspection and data cleaning.
- Check for missing data and formulate an apt strategy to treat them.
- Remove duplicate data records.
- Perform descriptive analytics on the given data.

2. Data Transformation:

- Perform cohort analysis (a cohort is a group of subjects that share a defining characteristic). Observe how a cohort behaves across time and compare it to other cohorts.
- Create month cohorts and analyze active customers for each cohort.
- Analyze the retention rate of customers.

3. Data Modeling :

- Build a RFM (Recency Frequency Monetary) model. Recency means the number of days since a customer made the last purchase. Frequency is the number of purchase in a given period. It could be 3 months, 6 months or 1 year. Monetary is the total amount of money a customer spent in that given period. Therefore, big spenders will be differentiated among other customers such as MVP (Minimum Viable Product) or VIP.
- Calculate RFM metrics.
- Build RFM Segments. Give recency, frequency, and monetary scores individually by dividing them into quartiles.
- Combine three ratings to get a RFM segment (as strings).
- Get the RFM score by adding up the three ratings.
- Analyze the RFM segments by summarizing them and comment on the findings.

Note:

- Rate “recency” for customer who has been active more recently higher than the less recent customer, because each company wants its customers to be recent.
- Rate “frequency” and “monetary” higher, because the company wants the customer to visit more often and spend more money.

Project Task: Week 2

4. Data Modeling :

- Create clusters using k-means clustering algorithm.
- Prepare the data for the algorithm. If the data is asymmetrically distributed, manage the skewness with appropriate transformation. Standardize the data.
- Decide the optimum number of clusters to be formed.
- Analyze these clusters and comment on the results.

5. Create a dashboard in tableau by choosing appropriate chart types and metrics useful for the business. The dashboard must entail the following:

- Country-wise analysis to demonstrate average spend. Use a bar chart to show the monthly figures
- Bar graph of top 15 products which are mostly ordered by the users to show the number of products sold
- Bar graph to show the count of orders vs. hours throughout the day
- Plot the distribution of RFM values using histogram and frequency charts
- Plot error (cost) vs. number of clusters selected
- Visualize to compare the RFM values of the clusters using heatmap

Importing some of the required libraries

```
In [211]: 1 import numpy as np
          2 import pandas as pd
          3 import matplotlib.pyplot as plt
          4 %matplotlib inline
          5 import seaborn as sns
          6 import warnings
          7 warnings.filterwarnings('ignore')
```

Loading the dataset

```
In [212]: 1 retail_data=pd.read_excel('Online Retail.xlsx')
          2 retail_data.head()
```

Out[212]:

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
0	536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER	6	2010-12-01 08:26:00	2.55	17850.0	United Kingdom
1	536365	71053	WHITE METAL LANTERN	6	2010-12-01 08:26:00	3.39	17850.0	United Kingdom
2	536365	84406B	CREAM CUPID HEARTS COAT HANGER	8	2010-12-01 08:26:00	2.75	17850.0	United Kingdom
3	536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6	2010-12-01 08:26:00	3.39	17850.0	United Kingdom
4	536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6	2010-12-01 08:26:00	3.39	17850.0	United Kingdom

```
In [213]: 1 retail_data.shape
```

Out[213]: (541909, 8)

The dataset consists of **541909** records with **8** features

1. Data Cleaning:

The dataset consists of 'InvoiceNo', 'StockCode', 'Description', 'Quantity', 'InvoiceDate', 'UnitPrice', 'CustomerID', and 'Country' features.Out of these columns we can remove 'Description' column because this column does not provide any contribution to our model.

```
In [214]: 1 df=retail_data.drop(columns='Description')
          2 df.head()
```

Out[214]:

	InvoiceNo	StockCode	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
0	536365	85123A	6	2010-12-01 08:26:00	2.55	17850.0	United Kingdom
1	536365	71053	6	2010-12-01 08:26:00	3.39	17850.0	United Kingdom
2	536365	84406B	8	2010-12-01 08:26:00	2.75	17850.0	United Kingdom
3	536365	84029G	6	2010-12-01 08:26:00	3.39	17850.0	United Kingdom
4	536365	84029E	6	2010-12-01 08:26:00	3.39	17850.0	United Kingdom


a. Checking for missing values in the dataset

```
In [215]: 1 df.isnull().sum()
```

Out[215]: InvoiceNo 0
StockCode 0
Quantity 0
InvoiceDate 0
UnitPrice 0
CustomerID 135080
Country 0
dtype: int64

```
In [216]: 1 plt.figure(figsize=(5,4))
2 sns.heatmap(df.isnull(),cmap="crest_r",yticklabels=False,cbar=False)
3 plt.title('Visualizing missing values in the Dataset\n', fontsize = 12)
4 plt.xticks(rotation='vertical')
5 plt.show()
```

Visualizing missing values in the Dataset



```
In [217]: 1 print('The Percentage of missing values in the dataset = {}'.format(round((df['CustomerID'].isnull().sum()/df.shape[0])*100,2)))
2
```

The Percentage of missing values in the dataset = 24.93%

The **CustomerID** feature has **135080** missing values in it. Since we are performing **customer segmentation** our analysis is mainly centered around Customers CustomerID feature hence replacing these records containing missing values with mean/median imputation will not be a good idea. So let's check for the possible ways to impute these missing values by comparing **InvoiceNo** feature with records having missing **CustomerID**.

```
In [218]: 1 unique_invoiceno=set(df[df['CustomerID'].isnull()]['InvoiceNo'])
2 unique_invoiceno
```

```
Out[218]: {540673,
540674,
540675,
540676,
540677,
540678,
540679,
540681,
'C544049',
540683,
540684,
540685,
540693,
540694,
540695,
540696,
548886,
548887,
540699,
540698}
```

```
In [219]: 1 df[df['InvoiceNo'].isin(unique_invoiceno) & ~df['CustomerID'].isnull()]
```

```
Out[219]:
```

InvoiceNo	StockCode	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
-----------	-----------	----------	-------------	-----------	------------	---------

On comparing customers InvoiceNo with CustomerID we did not find any other same InvoiceNo's from the customers having missing CustomerID's. So let's drop the records with missing CustomerID.

```
In [220]: 1 df.dropna(inplace=True)
```

```
In [221]: 1 df.isnull().sum().sum()
```

Out[221]: 0

b. Checking for duplicate records in the dataset

```
In [222]: 1 df.duplicated().sum()
```

Out[222]: 5227

There are **5227 duplicate records** in the dataset.Lets remove them!

```
In [223]: 1 df.drop_duplicates(inplace=True)
```

```
In [224]: 1 df.shape
```

Out[224]: (401602, 7)

```
In [225]: 1 df.dtypes
```

Out[225]: InvoiceNo object
StockCode object
Quantity int64
InvoiceDate datetime64[ns]
UnitPrice float64
CustomerID float64
Country object
dtype: object

Lets change CustomerID to object because CustomerID will not be numeric in most of the cases

```
In [226]: 1 df['CustomerID'] = df['CustomerID'].astype(str)
```

```
In [227]: 1 df['CustomerID'].dtype
```

Out[227]: dtype('O')

c. Descriptive Analysis

```
In [228]: 1 df.describe(datetime_is_numeric=True).T
```

Out[228]:

	count	mean	min	25%	50%	75%	max	std
Quantity	401602.0	12.182579	-80995.0	2.0	5.0	12.0	80995.0	250.283248
InvoiceDate	401602	2011-07-10 12:08:08.129839872	2010-12-01 08:26:00	2011-04-06 15:02:00	2011-07-29 15:40:00	2011-10-20 11:58:00	2011-12-09 12:50:00	NaN
UnitPrice	401602.0	3.474064	0.0	1.25	1.95	3.75	38970.0	69.764209

```
In [229]: 1 df.describe(include='O').T
```

Out[229]:

	count	unique	top	freq
InvoiceNo	401602	22190	576339	542
StockCode	401602	3684	85123A	2065
CustomerID	401602	4372	17841.0	7812
Country	401602	37	United Kingdom	356726

From the above **Descriptive statistics** we find that,

- The company has customers across 37 Countries and most of the customers are from United Kingdom.
- The company has a total of 4372 unique customers among them a customer with CustomerID 17841.0 has done the most purchases.
- There are 3684 unique StockCodes which means that there are 3684 unique products in total and the product with StockCode 85123A is the most frequently purchased product.
- Out of 401602 Invoices we have 22190 unique invoices which implies 22190 unique transactions has been done.
- From the UnitPrice feature we can notice that each product in the transaction costs an avg of 3.47 sterling.
- From the InvoiceDate we see that we have transactional data from 1st december 2010 to 9th december 2011
- we also notice negative values in the Quantity feature which indicates that some of the customers has returned the products.

2. Data Transformation

Cohort analysis

Assigning the cohorts and calculating the monthly offset

```
In [230]: 1 # Creating a function that will parse the date Time based cohort: 1 day of month
2 import datetime as dt
3 def get_month(x) :
4     return dt.datetime(x.year,x.month,1)
```

```
In [231]: 1 #Creating InvoiceMonth column from InvoiceDate
2 df['InvoiceMonth'] = df['InvoiceDate'].apply(get_month)
3 #Grouping by CustomerId and selecting the InvoiceMonth
4 grouping = df.groupby('CustomerID')['InvoiceMonth']
5 # Assigning a minimum InvoiceMonth to the dataset
6 df['CohortMonth'] = grouping.transform('min')
7 df.head()
```

```
Out[231]:
```

	InvoiceNo	StockCode	Quantity	InvoiceDate	UnitPrice	CustomerID	Country	InvoiceMonth	CohortMonth
0	536365	85123A	6	2010-12-01 08:26:00	2.55	17850.0	United Kingdom	2010-12-01	2010-12-01
1	536365	71053	6	2010-12-01 08:26:00	3.39	17850.0	United Kingdom	2010-12-01	2010-12-01
2	536365	84406B	8	2010-12-01 08:26:00	2.75	17850.0	United Kingdom	2010-12-01	2010-12-01
3	536365	84029G	6	2010-12-01 08:26:00	3.39	17850.0	United Kingdom	2010-12-01	2010-12-01
4	536365	84029E	6	2010-12-01 08:26:00	3.39	17850.0	United Kingdom	2010-12-01	2010-12-01

Calculating time offset in Month as Cohort Index

```
In [232]: 1 def get_month_int (df,column):
2     year = df[column].dt.year
3     month = df[column].dt.month
4     day = df[column].dt.day
5     return year,month,day
6
7 #Getting the integers for date parts from the 'InvoiceDay' column
8 invoice_year,invoice_month,_ = get_month_int(df,'InvoiceMonth')
9 #Getting the integers for date parts from the 'CohortDay' column
10 cohort_year,cohort_month,_ = get_month_int(df,'CohortMonth')
11 #Calculating the difference in years
12 year_diff = invoice_year - cohort_year
13 #Calculating the difference in months
14 month_diff = invoice_month - cohort_month
15
16 df['CohortIndex'] = year_diff * 12 + month_diff + 1
```

```
In [233]: 1 df.head()
```

```
Out[233]:
```

	InvoiceNo	StockCode	Quantity	InvoiceDate	UnitPrice	CustomerID	Country	InvoiceMonth	CohortMonth	CohortIndex
0	536365	85123A	6	2010-12-01 08:26:00	2.55	17850.0	United Kingdom	2010-12-01	2010-12-01	1
1	536365	71053	6	2010-12-01 08:26:00	3.39	17850.0	United Kingdom	2010-12-01	2010-12-01	1
2	536365	84406B	8	2010-12-01 08:26:00	2.75	17850.0	United Kingdom	2010-12-01	2010-12-01	1
3	536365	84029G	6	2010-12-01 08:26:00	3.39	17850.0	United Kingdom	2010-12-01	2010-12-01	1
4	536365	84029E	6	2010-12-01 08:26:00	3.39	17850.0	United Kingdom	2010-12-01	2010-12-01	1

```
In [234]: 1 #Counting the monthly active customers from each cohort
2 grouping = df.groupby(['CohortMonth','CohortIndex'])
3 #Counting number of unique CustomerId's falling in each group of CohortMonth and CohortIndex
4 cohort_data = grouping['CustomerID'].apply(pd.Series.nunique)
5 cohort_data = cohort_data.reset_index()
6 cohort_data.head()
```

```
Out[234]:
```

	CohortMonth	CohortIndex	CustomerID
0	2010-12-01	1	948
1	2010-12-01	2	362
2	2010-12-01	3	317
3	2010-12-01	4	367
4	2010-12-01	5	341

Calculating Retention rate

In [235]:

```
1 #Assigning column names to the dataframe created above
2 cohort_counts = cohort_data.pivot(index='CohortMonth',columns='CohortIndex',values='CustomerID')
3 cohort_counts
```

Out[235]:

CohortIndex	1	2	3	4	5	6	7	8	9	10	11	12	13
CohortMonth													
2010-12-01	948.0	362.0	317.0	367.0	341.0	376.0	360.0	336.0	336.0	374.0	354.0	474.0	260.0
2011-01-01	421.0	101.0	119.0	102.0	138.0	126.0	110.0	108.0	131.0	146.0	155.0	63.0	NaN
2011-02-01	380.0	94.0	73.0	106.0	102.0	94.0	97.0	107.0	98.0	119.0	35.0	NaN	NaN
2011-03-01	440.0	84.0	112.0	96.0	102.0	78.0	116.0	105.0	127.0	39.0	NaN	NaN	NaN
2011-04-01	299.0	68.0	66.0	63.0	62.0	71.0	69.0	78.0	25.0	NaN	NaN	NaN	NaN
2011-05-01	279.0	66.0	48.0	48.0	60.0	68.0	74.0	29.0	NaN	NaN	NaN	NaN	NaN
2011-06-01	235.0	49.0	44.0	64.0	58.0	79.0	24.0	NaN	NaN	NaN	NaN	NaN	NaN
2011-07-01	191.0	40.0	39.0	44.0	52.0	22.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2011-08-01	167.0	42.0	42.0	42.0	23.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2011-09-01	298.0	89.0	97.0	36.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2011-10-01	352.0	93.0	46.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2011-11-01	321.0	43.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2011-12-01	41.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

In [236]:

```
1 cohort_size = cohort_counts.iloc[:,0]
2 retention_rate = cohort_counts.divide(cohort_size,axis=0) # axis=0 to divide along the row axis
3 # Coverting the retention rate into percentage and Rounding off.
4 retention_rate.round(2) * 100
```

Out[236]:

CohortIndex	1	2	3	4	5	6	7	8	9	10	11	12	13
CohortMonth													
2010-12-01	100.0	38.0	33.0	39.0	36.0	40.0	38.0	35.0	35.0	39.0	37.0	50.0	27.0
2011-01-01	100.0	24.0	28.0	24.0	33.0	30.0	26.0	26.0	31.0	35.0	37.0	15.0	NaN
2011-02-01	100.0	25.0	19.0	28.0	27.0	25.0	26.0	28.0	26.0	31.0	9.0	NaN	NaN
2011-03-01	100.0	19.0	25.0	22.0	23.0	18.0	26.0	24.0	29.0	9.0	NaN	NaN	NaN
2011-04-01	100.0	23.0	22.0	21.0	21.0	24.0	23.0	26.0	8.0	NaN	NaN	NaN	NaN
2011-05-01	100.0	24.0	17.0	17.0	22.0	24.0	27.0	10.0	NaN	NaN	NaN	NaN	NaN
2011-06-01	100.0	21.0	19.0	27.0	25.0	34.0	10.0	NaN	NaN	NaN	NaN	NaN	NaN
2011-07-01	100.0	21.0	20.0	23.0	27.0	12.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2011-08-01	100.0	25.0	25.0	25.0	14.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2011-09-01	100.0	30.0	33.0	12.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2011-10-01	100.0	26.0	13.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2011-11-01	100.0	13.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2011-12-01	100.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

Customer retention is a very useful metric to understand how many of the customers are still active out of all the customers.Retention actually gives you the percentage of active customers compared to the total number of customers.

The above retention rate dataframe represents Customer retained across Cohorts.We can read it as follows:

- Index value represents the Cohort
- Columns represent the number of months since the current Cohort
- For instance - The value at CohortMonth 2011-01-01, CohortIndex 5 is 33.0 and represents 33% of customers from cohort 2011-01 were retained in the 5th Month.


Also, we can notice from the retention Rate DataFrame:

- Retention Rate 1st index i.e 1st month is 100% as all the customers for that particular customer signed up in 1st Month
- The retention rate may increase or decrease in subsequent Indexes.
- Values towards the bottom right have a lot of NaN values.

Visualizing the Retention rate

In [237]:

```
1 #Creating the heatmap
2 plt.figure(figsize=(14, 8))
3 sns.heatmap(retention_rate, annot = True,vmin = 0.0, vmax =0.5,cmap="crest_r", fmt='.0%')
4 plt.title('Retention Rate in Percentage : Monthly Cohorts\n', fontsize = 18)
5 plt.ylabel('Cohort Month',fontsize = 14)
6 plt.xlabel('Cohort Index',fontsize = 14)
7 plt.show()
```



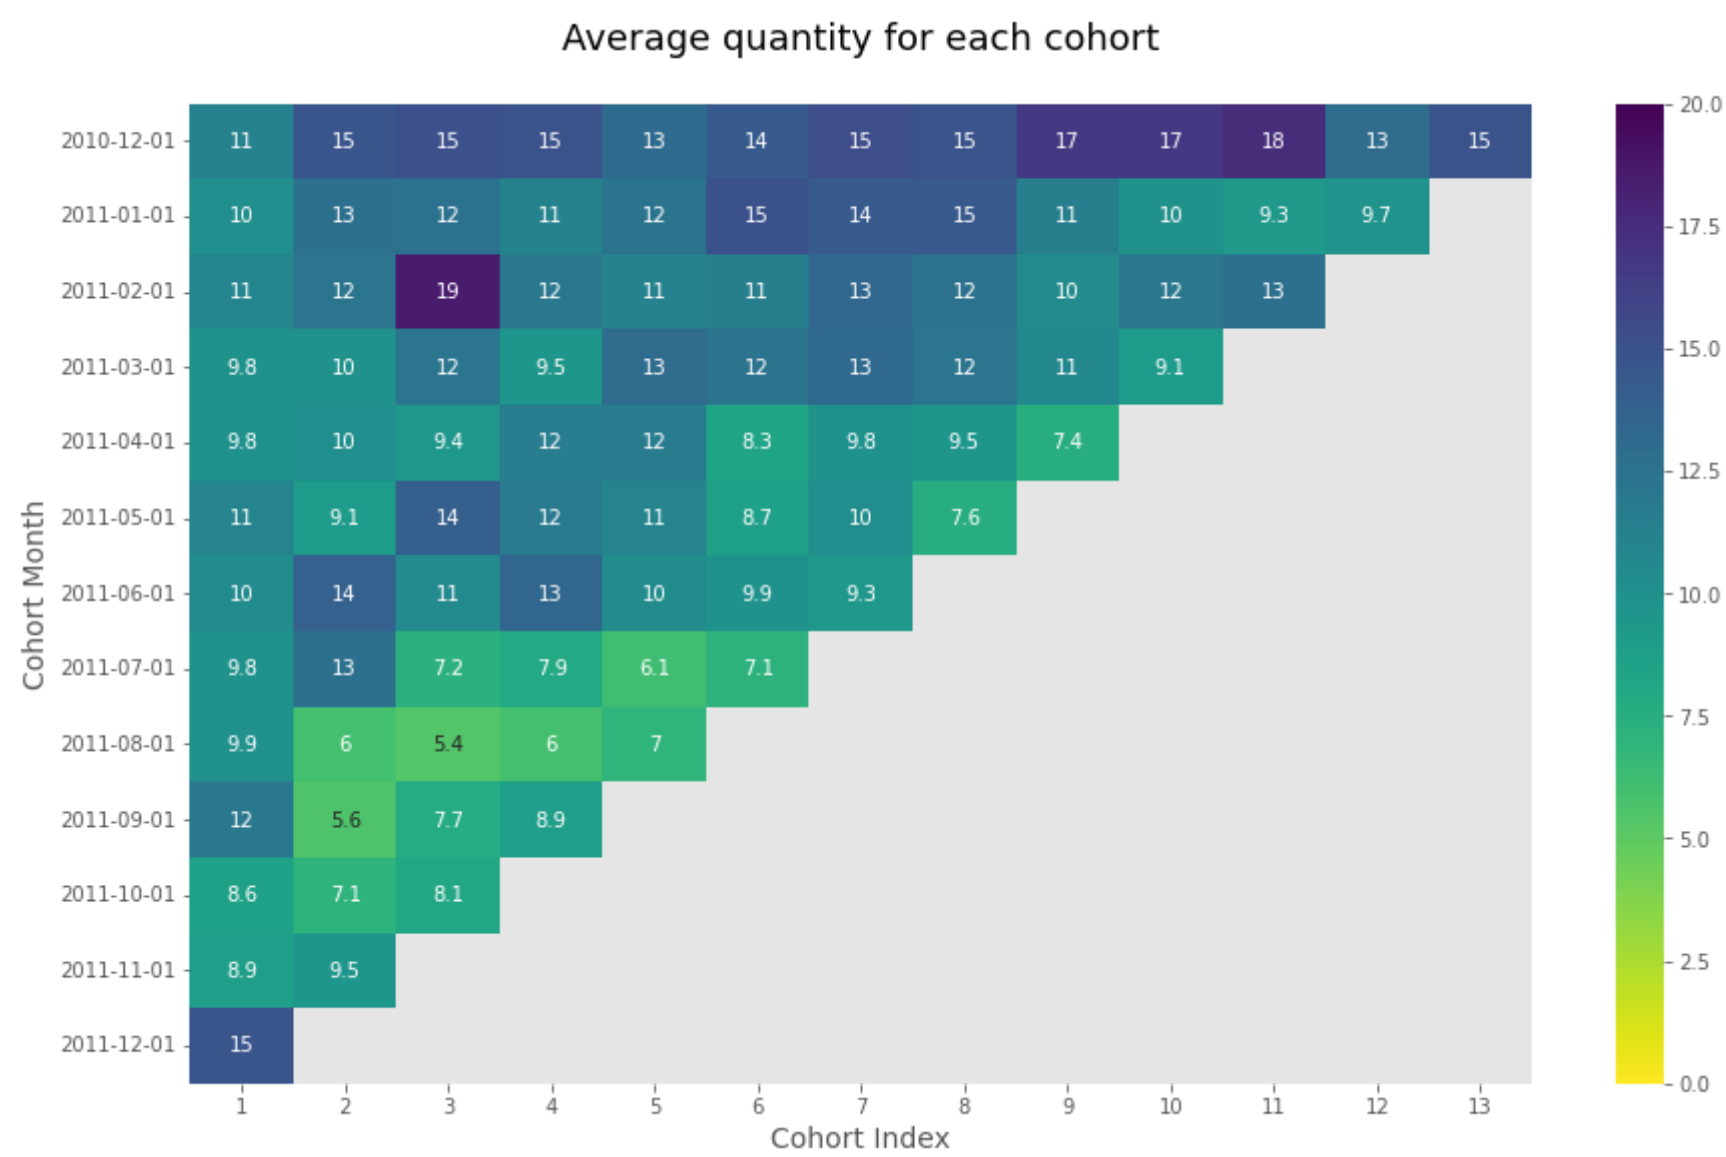
Average quantity for each cohort

In [238]:

```
1 grouping1 = df.groupby(['CohortMonth', 'CohortIndex'])
2 cohort_data1 = grouping1['Quantity'].mean()
3 cohort_data1 = cohort_data1.reset_index()
4 average_quantity = cohort_data1.pivot(index='CohortMonth',columns='CohortIndex',values='Quantity')
5 average_quantity.round(2)
6 average_quantity.index = average_quantity.index.date
```

In [239]:

```
1 #Creating the heatmap
2 plt.figure(figsize=(15, 9))
3 sns.heatmap(average_quantity, annot = True,vmin = 0.0, vmax =20,cmap="viridis_r")
4 plt.title('Average quantity for each cohort\n', fontsize = 18)
5 plt.ylabel('Cohort Month',fontsize = 14)
6 plt.xlabel('Cohort Index',fontsize = 14)
7 plt.show()
```



3. Data Modelling

a. Building a RFM (Recency Frequency Monetary) model

RFM

- **Recency** : The freshness of the customer activity be it purchases or visits (or) Recency means the number of days since a customer made the last purchase

- **Frequency** : The frequency of the customer transactions or visits (or) Frequency is the number of purchase in a given period. It could be 3 months, 6 months or 1 year.
- **Monetary** : The intension of customer to spend or purchasing power of customer (or) Monetary is the total amount of money a customer spent in that given period. Therefore, big spenders will be differentiated among other customers such as MVP (Minimum Viable Product) or VIP

The RFM values can be grouped in several ways:

1. Percentiles
2. Pareto (80/20)
3. Custom - based on business knowledge

Lets implement Percentile based grouping To calculate percentiles:

1. Sort customers based on that metric
2. Break customers into a pre-defined number of groups of equal size
3. Assign a label to each group

```
In [240]: 1 #Calculating TotalAmount
2 df['TotalAmount'] = df['UnitPrice']* df['Quantity']
3 df.head()
```

```
Out[240]:
```

	InvoiceNo	StockCode	Quantity	InvoiceDate	UnitPrice	CustomerID	Country	InvoiceMonth	CohortMonth	CohortIndex	TotalAmount
0	536365	85123A	6	2010-12-01 08:26:00	2.55	17850.0	United Kingdom	2010-12-01	2010-12-01	1	15.30
1	536365	71053	6	2010-12-01 08:26:00	3.39	17850.0	United Kingdom	2010-12-01	2010-12-01	1	20.34
2	536365	84406B	8	2010-12-01 08:26:00	2.75	17850.0	United Kingdom	2010-12-01	2010-12-01	1	22.00
3	536365	84029G	6	2010-12-01 08:26:00	3.39	17850.0	United Kingdom	2010-12-01	2010-12-01	1	20.34
4	536365	84029E	6	2010-12-01 08:26:00	3.39	17850.0	United Kingdom	2010-12-01	2010-12-01	1	20.34

```
In [241]: 1 print('Min Invoice Date:{}\nMax Invoice Date:{}'.format(df.InvoiceDate.dt.date.min(),df.InvoiceDate.dt.date.max()))

Min Invoice Date:2010-12-01
Max Invoice Date:2011-12-09
```

```
In [242]: 1 #In the real world, we will be working with the most recent snapshot of the data of today or yesterday
2 # so Lets count different days with snapshot_date.
3 snapshot_date = df['InvoiceDate'].max() + dt.timedelta(days=1)
4 snapshot_date
```

```
Out[242]: Timestamp('2011-12-10 12:50:00')
```

b. Calculating RFM metrics

```
In [243]: 1 #lambda function used below gives the number of days between hypothetical today and the last transaction
2 rfm = df.groupby(['CustomerID']).agg({'InvoiceDate': lambda x : (snapshot_date - x.max()).days,
3                                     'InvoiceNo': 'count', 'TotalAmount': 'sum'})
4 #lets rename the columns
5 rfm.rename(columns={'InvoiceDate': 'Recency', 'InvoiceNo': 'Frequency', 'TotalAmount': 'Monetary'})
6         ,inplace= True)
7
8 #Final RFM values
9 rfm.head()
```

```
Out[243]:
```

CustomerID	Recency	Frequency	Monetary
12346.0	326	2	0.00
12347.0	2	182	4310.00
12348.0	75	31	1797.24
12349.0	19	73	1757.55
12350.0	310	17	334.40

NOTE :

- We will rate "Recency" customers who have been active more recently better than the less recent customer,because each company wants its customers to be recent.
- We will also rate "Frequency" and "Monetary" higher label because we want Customer to spend more money and visit more often.

c. Building RFM Segments, and producing recency, frequency, and monetary scores individually by

dividing them into quartiles

```
In [246]: 1 #Building RFM segments
2 r_labels =range(4,0,-1)
3 f_labels=range(1,5)
4 m_labels=range(1,5)
5 r_quartiles = pd.qcut(rfm['Recency'], q=4, labels = r_labels)
6 f_quartiles = pd.qcut(rfm['Frequency'],q=4, labels = f_labels)
7 m_quartiles = pd.qcut(rfm['Monetary'],q=4,labels = m_labels)
8 rfm = rfm.assign(R=r_quartiles,F=f_quartiles,M=m_quartiles)
9
10 #Combining three ratings to get a RFM segment (as strings)
11 def add_rfm(x) : return str(x['R']) + str(x['F']) + str(x['M'])
12 rfm['RFM_Segment'] = rfm.apply(add_rfm,axis=1 )
13
14 #Getting the RFM score by adding up the three ratings
15 rfm['RFM_Score'] = rfm[['R','F','M']].sum(axis=1)
16
17 #displaying the dataframe which has a row for each customer with their RFM
18 rfm.head()
```

```
Out[246]:
```

	Recency	Frequency	Monetary	R	F	M	RFM_Segment	RFM_Score
CustomerID								
12346.0	326	2	0.00	1	1	1	111	3
12347.0	2	182	4310.00	4	4	4	444	12
12348.0	75	31	1797.24	2	2	4	224	8
12349.0	19	73	1757.55	3	3	4	334	10
12350.0	310	17	334.40	1	1	2	112	4

c. Analyzing the RFM segments

```
In [247]: 1 #It is always the best practice to investigate the size(largest RFM segments) of the segments
2 #before we use them for targeting or for other business Applications.
3 rfm.groupby(['RFM_Segment']).size().sort_values(ascending=False)[:5]
```

```
Out[247]: RFM_Segment
444      470
111      393
344      210
122      204
211      181
dtype: int64
```

Summary metrics per RFM Score

```
In [248]: 1 rfm.groupby('RFM_Score').agg({'Recency': 'mean', 'Frequency': 'mean',
2                                         'Monetary': ['mean', 'count'] }).round(2)
```

```
Out[248]:
```

	Recency	Frequency	Monetary	
	mean	mean	mean	count
RFM_Score				
3	264.75	7.85	109.25	393
4	174.88	13.79	226.57	390
5	152.50	20.92	342.69	515
6	94.59	28.11	490.78	469
7	79.49	39.22	722.68	439
8	62.76	56.08	965.48	467
9	44.37	77.38	1341.58	413
10	31.52	112.60	1819.18	442
11	20.47	190.42	3890.65	374
12	6.71	367.72	8848.13	470

Using RFM score lets group the customers into different segements like Gold, Silver and Bronze segments

```
In [249]: 1 def segments(df):
2         if df['RFM_Score'] > 9 :
3             return 'Gold'
4         elif (df['RFM_Score'] > 5) and (df['RFM_Score'] <= 9 ):
5             return 'Silver'
6         else:
7             return 'Bronze'
```


```
In [250]: 1 rfm['Customer_Segment'] = rfm.apply(segments,axis=1)
2         rfm.groupby('Customer_Segment').agg({'Recency':'mean', 'Frequency':'mean',
3                                               'Monetary': ['mean', 'count']}).round(2)
```

Out[250]:

	Recency	Frequency	Monetary	
	mean	mean	mean	count
Customer_Segment				
Bronze	193.21	14.82	237.12	1298
Gold	19.24	228.47	4990.52	1286
Silver	70.97	49.52	868.22	1788

```
In [251]: 1 plt.figure(figsize=(12,6))
2         plt.pie(rfm['Customer_Segment'].value_counts(),labels=rfm['Customer_Segment'].value_counts().index,
3               autopct='%.0f%',colors=['silver','crimson','gold'])
4         plt.title('Pie chart showing the distribution of customer segments',fontsize=14)
5         plt.show()
```

Pie chart showing the distribution of customer segments



d. Create clusters using K-means clustering algorithm

Preparing the data to implement algorithm

Lets check some of the assumptions before we implement our Kmeans Clustering algorithm

- Symmetric distribution of variables (not skewed)
- Variables with same mean
- Variables with same variance

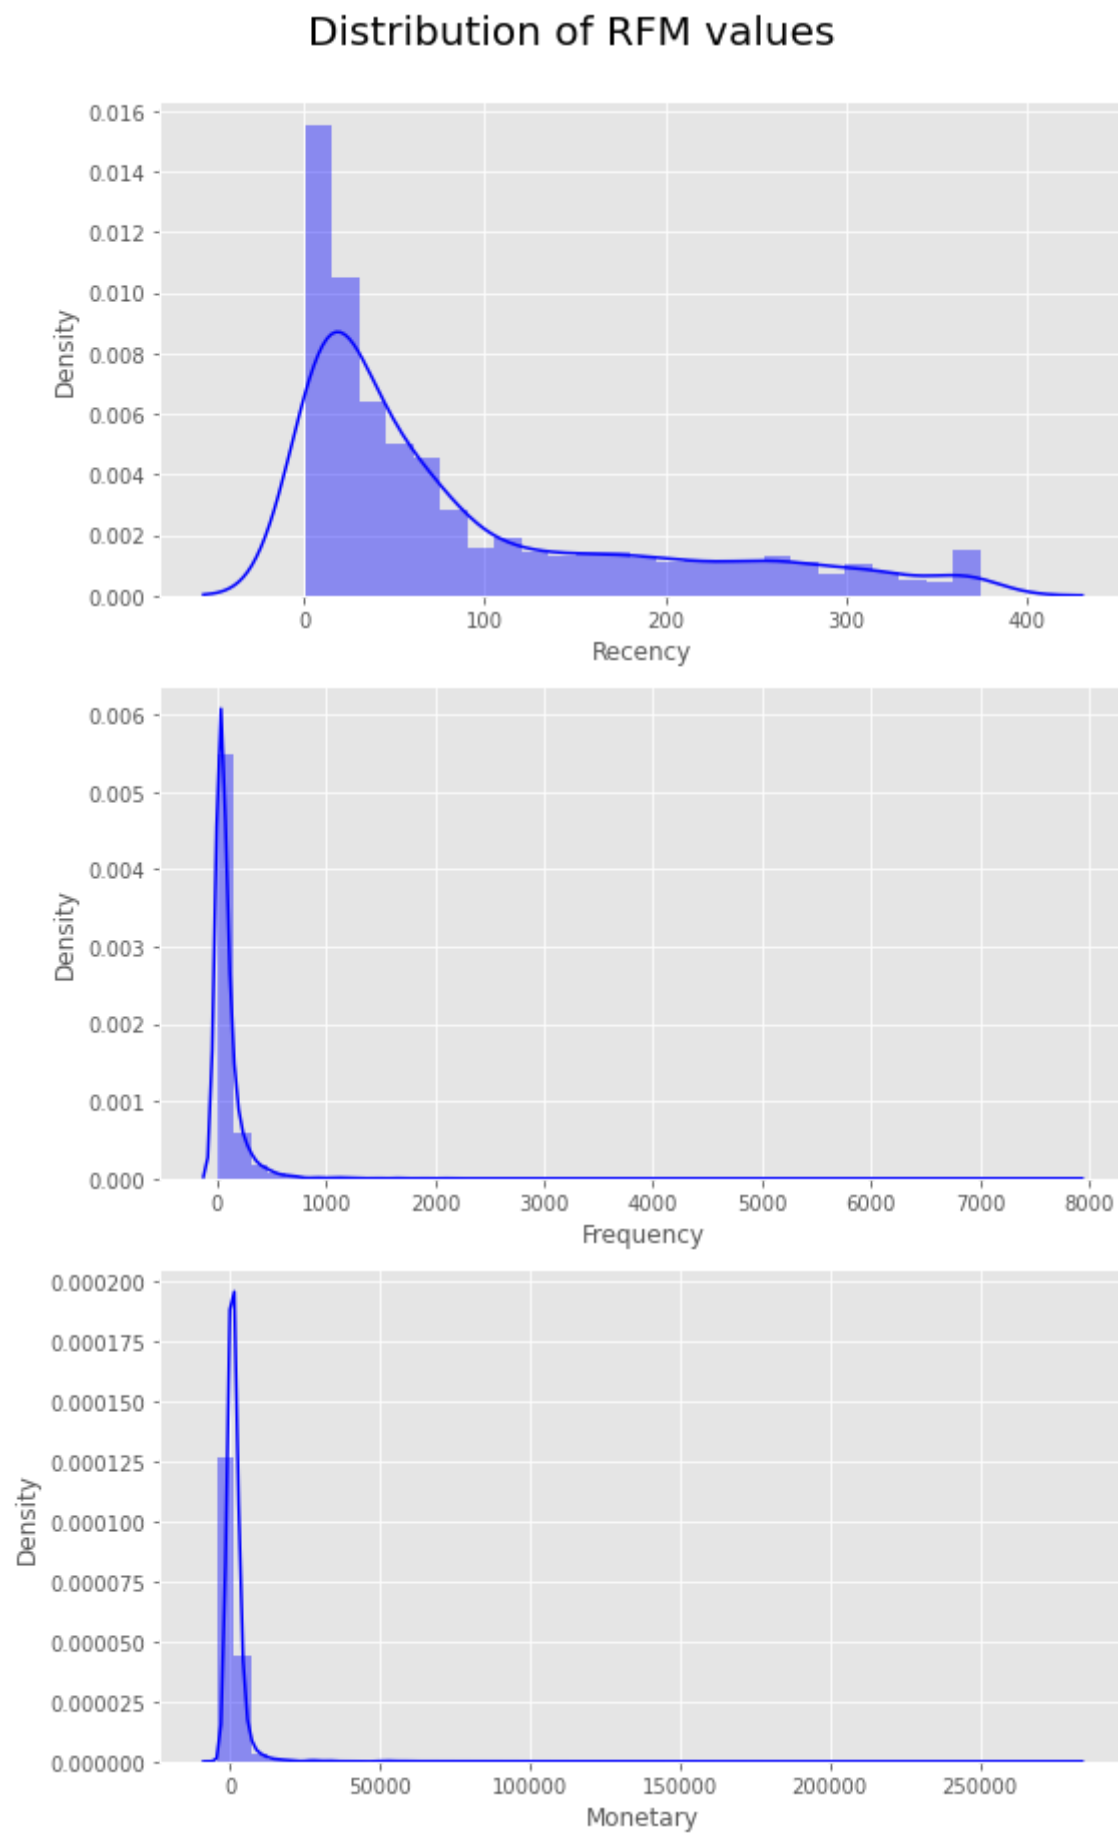
```
In [252]: 1 RFM=rfm[['Recency','Frequency','Monetary']]
2         RFM.describe().T
```

Out[252]:

	count	mean	std	min	25%	50%	75%	max
Recency	4372.0	92.047118	100.765435	1.00	17.000	50.00	143.000	374.00
Frequency	4372.0	91.857731	229.223553	1.00	17.000	41.00	99.250	7812.00
Monetary	4372.0	1893.405209	8218.228204	-4287.63	291.795	644.07	1608.335	279489.02

From the above RFM Descriptive statistics we observe that the **mean and variance are in different scales**,hence to overcome this problem we shall apply Scaling for variables to standardize the data.

```
In [253]: 1 # plot the distribution of RFM values
2 f,ax = plt.subplots(figsize=(8,13))
3 plt.suptitle('Distribution of RFM values\n',fontsize=20)
4 plt.subplot(311)
5 sns.distplot(RFM.Recency, label = 'Recency',color='b')
6 plt.subplot(312)
7 sns.distplot(RFM.Frequency, label = 'Frequency',color='b')
8 plt.subplot(313)
9 sns.distplot(RFM.Monetary, label = 'Monetary',color='b')
10 plt.style.context('fivethirtyeight')
11 plt.tight_layout()
12 plt.show()
```



From the above distribution plots we see that the variables are not Symmetrically distributed(i.e., data is skewed) hence we have to manage the skewness with appropriate transformation.

Lets Apply log Transformation to unskew the variables

In [254]:

```
1 rfm_log = RFM[['Recency', 'Frequency']].apply(np.log, axis = 1).round(2)
2 # Note: we can't take the Log of -ve numbers.
3 #from the descriptive statistics we found that 'Monetary' variable has -ve values in it
4 #hence we do some mathematical changes to make them +ve
5 rfm_log['Monetary'] = (RFM['Monetary']-RFM['Monetary'].min()+1).apply(np.log, axis = 1).round(2)
```

In [255]:

```
1 rfm_log.head()
```

Out[255]:

	Recency	Frequency	Monetary
CustomerID			
12346.0	5.79	0.69	8.36
12347.0	0.69	5.20	9.06
12348.0	4.32	3.43	8.71
12349.0	2.94	4.29	8.71
12350.0	5.74	2.83	8.44

In [256]:

```
1 rfm_log.describe().T
```

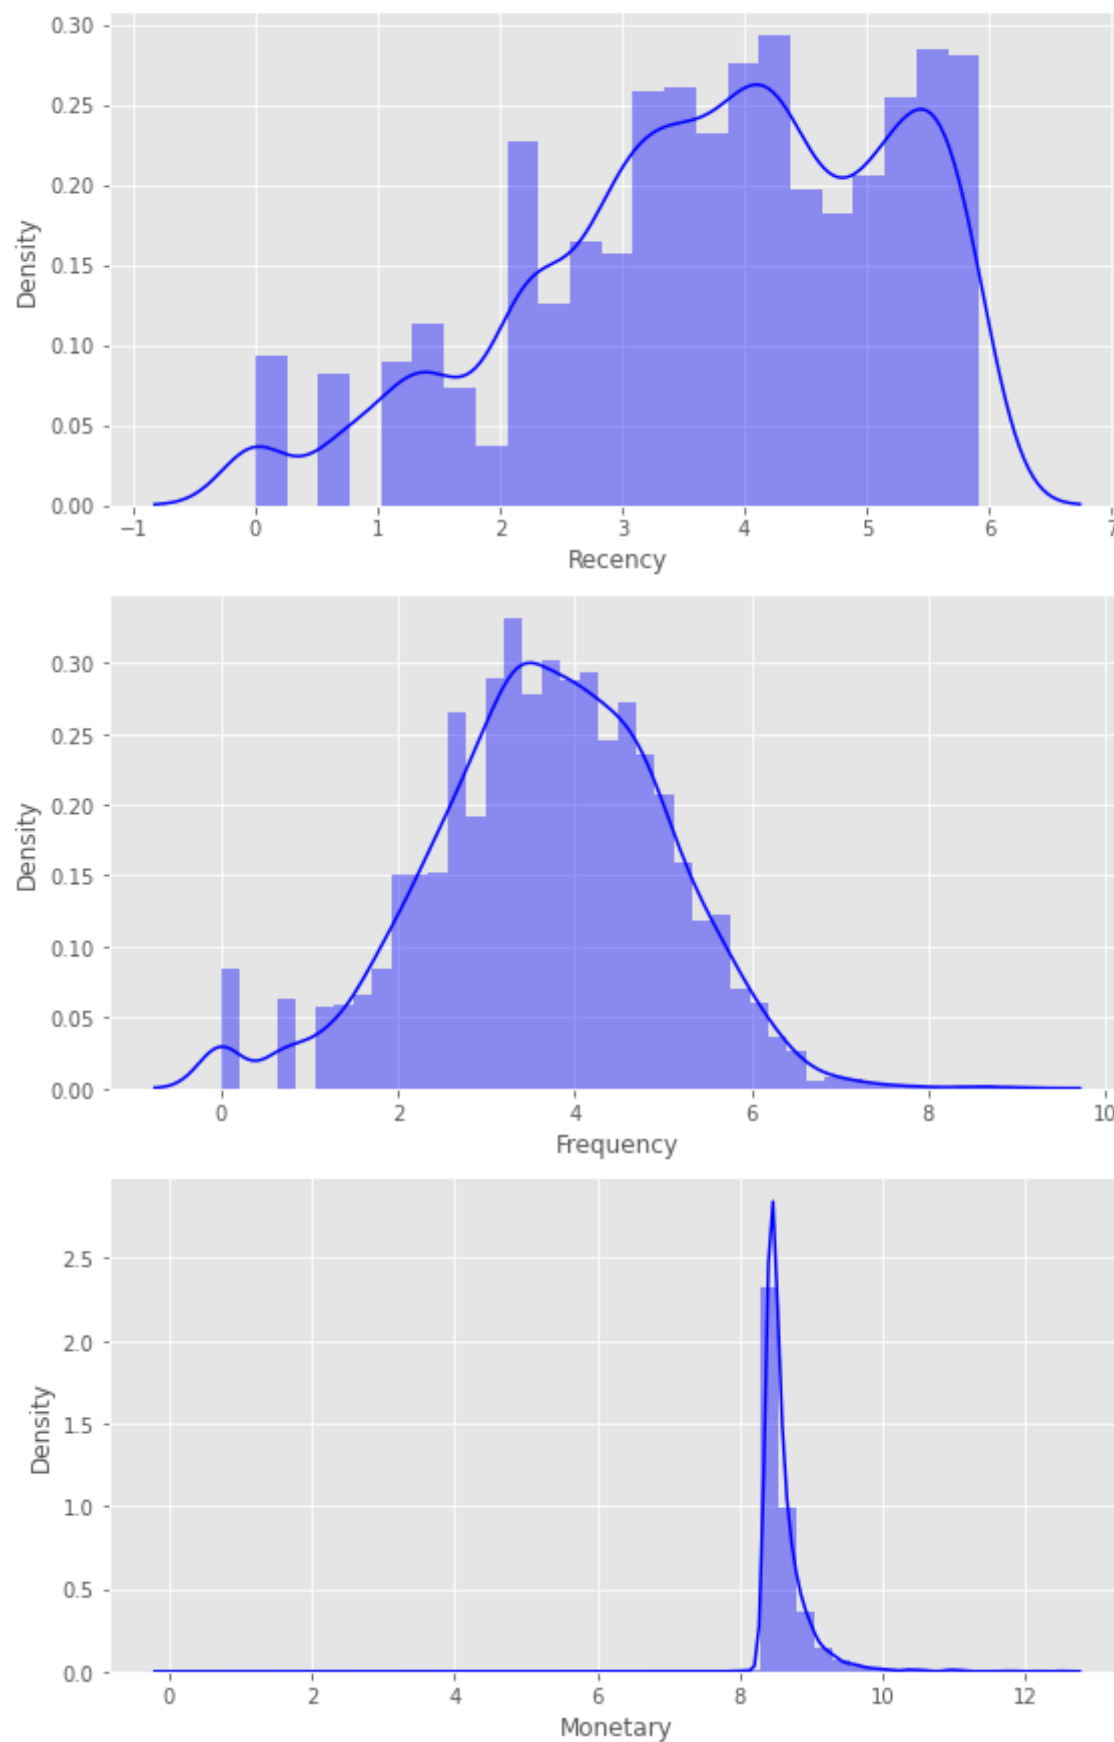
Out[256]:

	count	mean	std	min	25%	50%	75%	max
Recency	4372.0	3.732223	1.464554	0.0	2.83	3.91	4.9600	5.92
Frequency	4372.0	3.675672	1.335948	0.0	2.83	3.71	4.6025	8.96
Monetary	4372.0	8.615762	0.363824	0.0	8.43	8.50	8.6800	12.56

In [257]:

```
1 # plot the distribution of RFM values
2 f,ax = plt.subplots(figsize=(8,13))
3 plt.suptitle('Distribution of RFM values after Transformation\n',fontsize=20)
4 plt.subplot(311)
5 sns.distplot(rfm_log.Recency, label = 'Recency',color='b')
6 plt.subplot(312)
7 sns.distplot(rfm_log.Frequency, label = 'Frequency',color='b')
8 plt.subplot(313)
9 sns.distplot(rfm_log.Monetary, label = 'Monetary',color='b')
10 plt.style.context('fivethirtyeight')
11 plt.tight_layout()
12 plt.show()
```

Distribution of RFM values after Transformation



Implementation of K-Means Clustering Algorithm

Data PreProcessing


In [258]:

```
1 #Standardizing the variables with StandardScaler
2 from sklearn.preprocessing import StandardScaler
3 scaler = StandardScaler()
4 scaler.fit(rfm_log)
5 rfm_normalized= scaler.transform(rfm_log)
```

Choosing the Number of Clusters

Using Elbow method :

```
In [259]: 1 from sklearn.cluster import KMeans
2 # Finding the Optimal Number of Clusters with the help of Elbow Curve
3 # First : Get the Best KMeans
4 range_n_clusters = range(2,11)
5 inertias=[]
6 for k in range_n_clusters :
7     # Create a KMeans clusters
8     kc = KMeans(n_clusters=k,random_state=1,max_iter=50)
9     kc.fit(rfm_normalized)
10    inertias.append(kc.inertia_)
11
12 # Plot ks vs inertias
13 f, ax = plt.subplots(figsize=(8,4))
14 plt.plot(range_n_clusters , inertias, '-o',color='b')
15 plt.xlabel('Number of clusters, k')
16 plt.ylabel('Inertia')
17 plt.xticks(range_n_clusters)
18 plt.title('What is the best number of clusters for Kmeans using Elbow method?\n',fontsize=14)
19 plt.show()
```



In the above plot y axis represents **Inertia** which is the sum of squared distances of samples to their closest cluster centre and x axis represents the Number of cluster. Based on the observation,the **k-value of 3** is the best hyperparameter for our model because the next k-value tend to have a linear trend.

```
In [260]: 1 # Creating a dataframe for exporting to create visualization in tableau later!
2 df_inertia = pd.DataFrame(list(zip(range_n_clusters, inertias)), columns=['clusters', 'intertia'])
3 df_inertia
```

Out[260]:

	clusters	intertia
0	2	7829.819111
1	3	6056.792083
2	4	5055.953421
3	5	4266.736276
4	6	3720.831491
5	7	3185.333092
6	8	2755.271136
7	9	2498.757205
8	10	2279.985651

Using Silhouette Analysis :

```
In [261]: 1 # Lets also find the Optimal Number of Clusters with the help of Silhouette Analysis
2 from sklearn.metrics import silhouette_score
3 range_n_clusters = range(3,11)
4 for num_clusters in range_n_clusters:
5     kmeans = KMeans(n_clusters=num_clusters,random_state=1,max_iter=50)
6     kmeans.fit(rfm_normalized)
7     cluster_labels = kmeans.labels_
8     silhouette_avg = silhouette_score(rfm_normalized, cluster_labels)
9     print("For {0} clusters,the silhouette score is {1}".format(num_clusters, silhouette_avg))
```

For 3 clusters,the silhouette score is 0.3481717159163553
For 4 clusters,the silhouette score is 0.3015186069674956
For 5 clusters,the silhouette score is 0.2997450778586549
For 6 clusters,the silhouette score is 0.29354665467502505
For 7 clusters,the silhouette score is 0.295060959154092
For 8 clusters,the silhouette score is 0.3067768527176323
For 9 clusters,the silhouette score is 0.29667677045214796
For 10 clusters,the silhouette score is 0.29656278801153885

From the above analysis we see that the silhouette score is maximized at k = 3 so from this analysis also we find that selecting **3 as the optimum number of clusters** will be better for our model!

Applying Kmeans Clustering

```
In [262]: 1 # clustering
2 kc = KMeans(n_clusters= 3,random_state=1,max_iter=50)
3 kc.fit(rfm_normalized)
4
5 #Create a cluster label column in the original DataFrame
6 cluster_labels = kc.labels_
7
8 #Calculating the average RFM values and size for each cluster:
9 rfm_k3= RFM.assign(K_Cluster = cluster_labels)
10
11 #Calculating the average RFM values and sizes for each cluster:
12 rfm_k3.groupby('K_Cluster').agg({'Recency': 'mean','Frequency': 'mean',
13                                'Monetary': ['mean', 'count'],}).round(2)
```

Out[262]:

	Recency	Frequency	Monetary	
	mean	mean	mean	count
K_Cluster				
0	163.95	21.71	425.20	2025
1	13.08	409.90	11370.38	426
2	33.76	95.28	1339.48	1921

```
In [263]: 1 rfm['K_Cluster']=kc.labels_
2 rfm.head()
```

Out[263]:

	Recency	Frequency	Monetary	R	F	M	RFM_Segment	RFM_Score	Customer_Segment	K_Cluster
CustomerID										
12346.0	326	2	0.00	1	1	1	111	3	Bronze	0
12347.0	2	182	4310.00	4	4	4	444	12	Gold	1
12348.0	75	31	1797.24	2	2	4	224	8	Silver	0
12349.0	19	73	1757.55	3	3	4	334	10	Gold	2
12350.0	310	17	334.40	1	1	2	112	4	Bronze	0

Statistical Summary for RFM Quantiles

```
In [264]: 1 rfm.groupby('Customer_Segment').agg({'Recency':['mean','min','max'],
2                                           'Frequency':['mean','min','max'],
3                                           'Monetary':['mean','min','max']})
```

Out[264]:

	Recency		Frequency		Monetary	
	mean	min	max	mean	min	max
Customer_Segment						
Bronze	193.211094	18	374	14.821263	1	84
Gold	19.241835	1	140	228.471229	20	7812
Silver	70.971477	1	374	49.524609	1	526

Statistical Summary for Kmeans


```
In [265]: 1 rfm.groupby('K_Cluster').agg({'Recency':['mean','min','max'],
2                                     'Frequency':['mean','min','max'],
3                                     'Monetary':['mean','min','max']})
```

Out[265]:

	Recency			Frequency			Monetary		
	mean	min	max	mean	min	max	mean	min	max
K_Cluster									
0	163.949136	3	374	21.707160	1	155	425.204327	-4.287630e+03	7092.06
1	13.082160	1	267	409.896714	9	7812	11370.375188	1.863760e+03	279489.02
2	33.763665	1	315	95.277980	2	615	1339.484113	1.776357e-15	6977.04

```
In [266]: 1 rfm_normalized = pd.DataFrame(rfm_normalized,index=RFM.index,columns=RFM.columns)
2 rfm_normalized['K_Cluster'] = kc.labels_
3 rfm_normalized['Customer_Segment'] = rfm['Customer_Segment']
4 rfm_normalized.reset_index(inplace = True)
5
6 #Melt the data into a long format so RFM values and metric names are stored in 1 column each
7 rfm_melt = pd.melt(rfm_normalized,id_vars=['CustomerID','Customer_Segment','K_Cluster'],
8                   value_vars=['Recency', 'Frequency', 'Monetary'],
9                   var_name='Metric',value_name='Value')
10 rfm_melt.head()
```


Out[266]:

	CustomerID	Customer_Segment	K_Cluster	Metric	Value
0	12346.0	Bronze	0	Recency	1.405214
1	12347.0	Gold	1	Recency	-2.077473
2	12348.0	Silver	0	Recency	0.401381
3	12349.0	Gold	2	Recency	-0.540993
4	12350.0	Bronze	0	Recency	1.371070

Visualization using Snake Plots

Snake Plots describes the summarized table visually

```
In [267]: 1 f, (ax1, ax2) = plt.subplots(1,2, figsize=(14,5))
2 # Snake plots with RFM Quantiles
3 sns.lineplot(x = 'Metric', y = 'Value',hue = 'Customer_Segment',data = rfm_melt,ax=ax1)
4 # Snake plots with K-Means
5 sns.lineplot(x = 'Metric', y = 'Value',hue = 'K_Cluster', data = rfm_melt,ax=ax2)
6 plt.suptitle("Snake Plot of RFM",fontsize=15)
7 plt.show()
```



On Comparing the Snakeplots based on RFM Quantiles and Kmeans we observe that both the plots almost looks similar to each other,Which means both the segmentations has lead to similar kind of results.

Relative importance of segment attributes

- Useful technique to identify relative importance of each segment's attribute
- Calculate average values of each cluster
- Calculate average values of population
- Calculate importance score by dividing them and subtracting 1 (ensures 0 is returned when cluster average equals population average)

Let’s also try with heatmaps. Heat maps are a graphical representation of data where larger values were colored in darker scales and smaller values in lighter scales. We can compare the variance between the groups quite intuitively by colors.

```
In [268]: 1 # the farther the ratio is from 0, the more important that attribute is for a segment relative to the total population
2 cluster_avg = rfm_k3.groupby(['K_Cluster']).mean()
3 population_avg = RFM.mean()
4 relative_imp = cluster_avg / population_avg - 1
5 relative_imp.round(2)
```

Out[268]:


	Recency	Frequency	Monetary
K_Cluster			
0	0.78	-0.76	-0.78
1	-0.86	3.46	5.01
2	-0.63	0.04	-0.29

```
In [270]: 1 #the mean value in total
2 total_avg = rfm.iloc[:, 0:3].mean()
3 # calculating the proportional gap with total mean
4 cluster_avg = rfm.groupby('Customer_Segment').mean().iloc[:, 0:3]
5 prop_rfm = cluster_avg/total_avg - 1
6 prop_rfm.round(2)
```

Out[270]:

	Recency	Frequency	Monetary
Customer_Segment			
Bronze	1.10	-0.84	-0.87
Gold	-0.79	1.49	1.64
Silver	-0.23	-0.46	-0.54

```
In [271]: 1 # heatmap of RFM Quantiles
2 f, (ax1, ax2) = plt.subplots(1,2, figsize=(12,5))
3 sns.heatmap(data=prop_rfm, cmap= 'crest_r',annot = True,ax=ax1)
4 ax1.set(title = "Heatmap of RFM quantile")
5 # heatmap with Kmeans
6 sns.heatmap(data=relative_imp, annot=True, cmap='BuGn_r', ax=ax2)
7 ax2.set(title = "Heatmap of K-Means")
8 plt.suptitle("Heat Map of RFM\n",fontsize=20)
9 plt.show()
```



Conclusion :

We made two kinds of segmentation with RFM quantiles and K-Means clustering methods.

By Observing the above Snake plots and Heatmaps,we got to know how each segment differ from each other.

- We infer that Customers in **Cluster 0 (or) Bronze segment** are less frequent buyers,spending low amount and also they have not purchased anything in recent times, hence they will be considered as **least important customers**.
- Then the Customers in **Cluster 1 (or) Gold segment** are the most frequent buyers, spending high amount and are also placed orders recently so they will be considered as the **most important customers**.
- Finally the Customers in **Cluster 2 (or) Silver segment** are the customers having Recency, Frequency, and Monetary value in the medium range ,hence they will also be considered as **important customers**.

4. Data Reporting

1. Create a dashboard in tableau by choosing appropriate chart types and metrics useful for the business. The dashboard must entail the following:

1. Country-wise analysis to demonstrate average spend. Use a bar chart to show the monthly figures

- 2. Bar graph of top 15 products which are mostly ordered by the users to show the number of products sold
- 3. Bar graph to show the count of orders vs. hours throughout the day
- 4. Plot the distribution of RFM values using histogram and frequency charts
- 5. Plot error (cost) vs. number of clusters selected
- 6. Visualize to compare the RFM values of the clusters using heatmap


Saving the Data into excel to use it in Tableau to create a Dashboard

```
In [272]: 1 writer = pd.ExcelWriter(r'C:\Users\Sudhakar\Simplilearn Capstone Project\Project 3\RetailOutputData.xlsx',
2               engine='xlsxwriter')
3 df.to_excel(writer, sheet_name='Retail_data', index=False)
4 rfm.to_excel(writer, sheet_name='rfm_data', index=False)
5 rfm_melt.to_excel(writer, sheet_name='rfm_melt', index=False)
6 df_inertia.to_excel(writer, sheet_name='inertia', index=False)
7 writer.save()
```

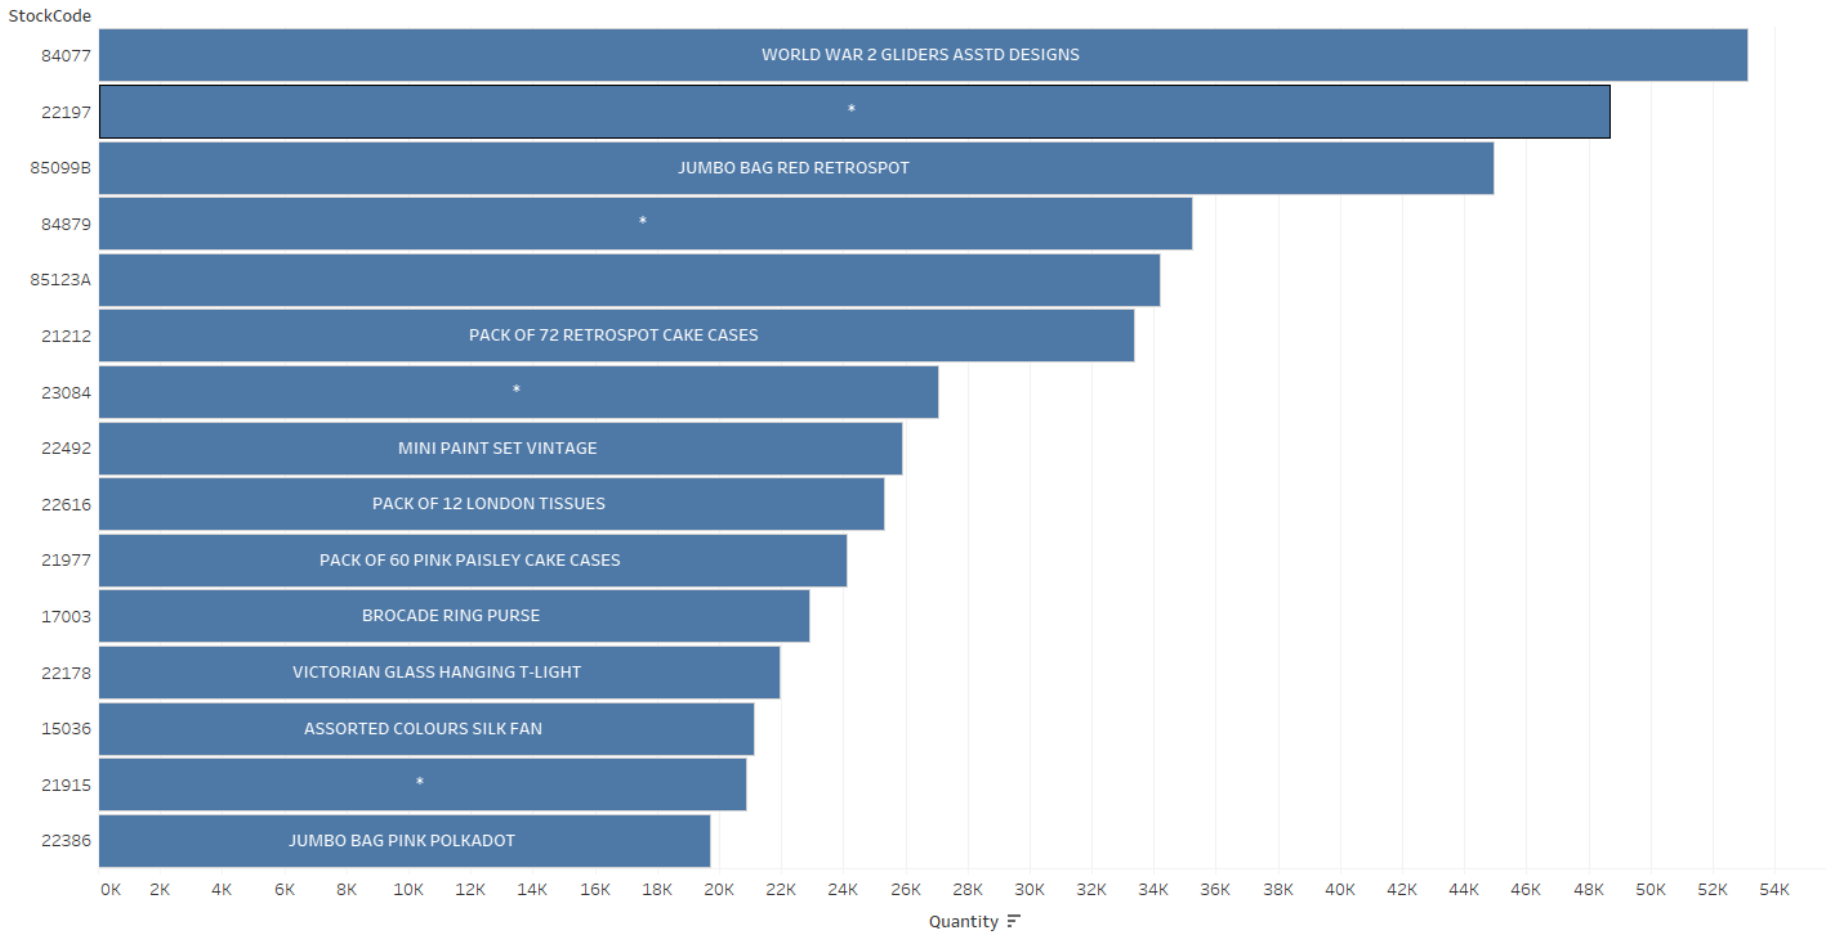
Please refer the below link for Data Reporting :
<https://public.tableau.com/app/profile/divya.reddy7718/viz/RetailCapstoneProjectDataReporting/Dashboard1>
(<https://public.tableau.com/app/profile/divya.reddy7718/viz/RetailCapstoneProjectDataReporting/Dashboard1>)

Screenshots of worksheets from Tableau


Country-Wise Monthly Average Spend




Top 15 products which are mostly ordered by the users




Count of orders Vs. Hours throughout the day




Distribution of RFM Values



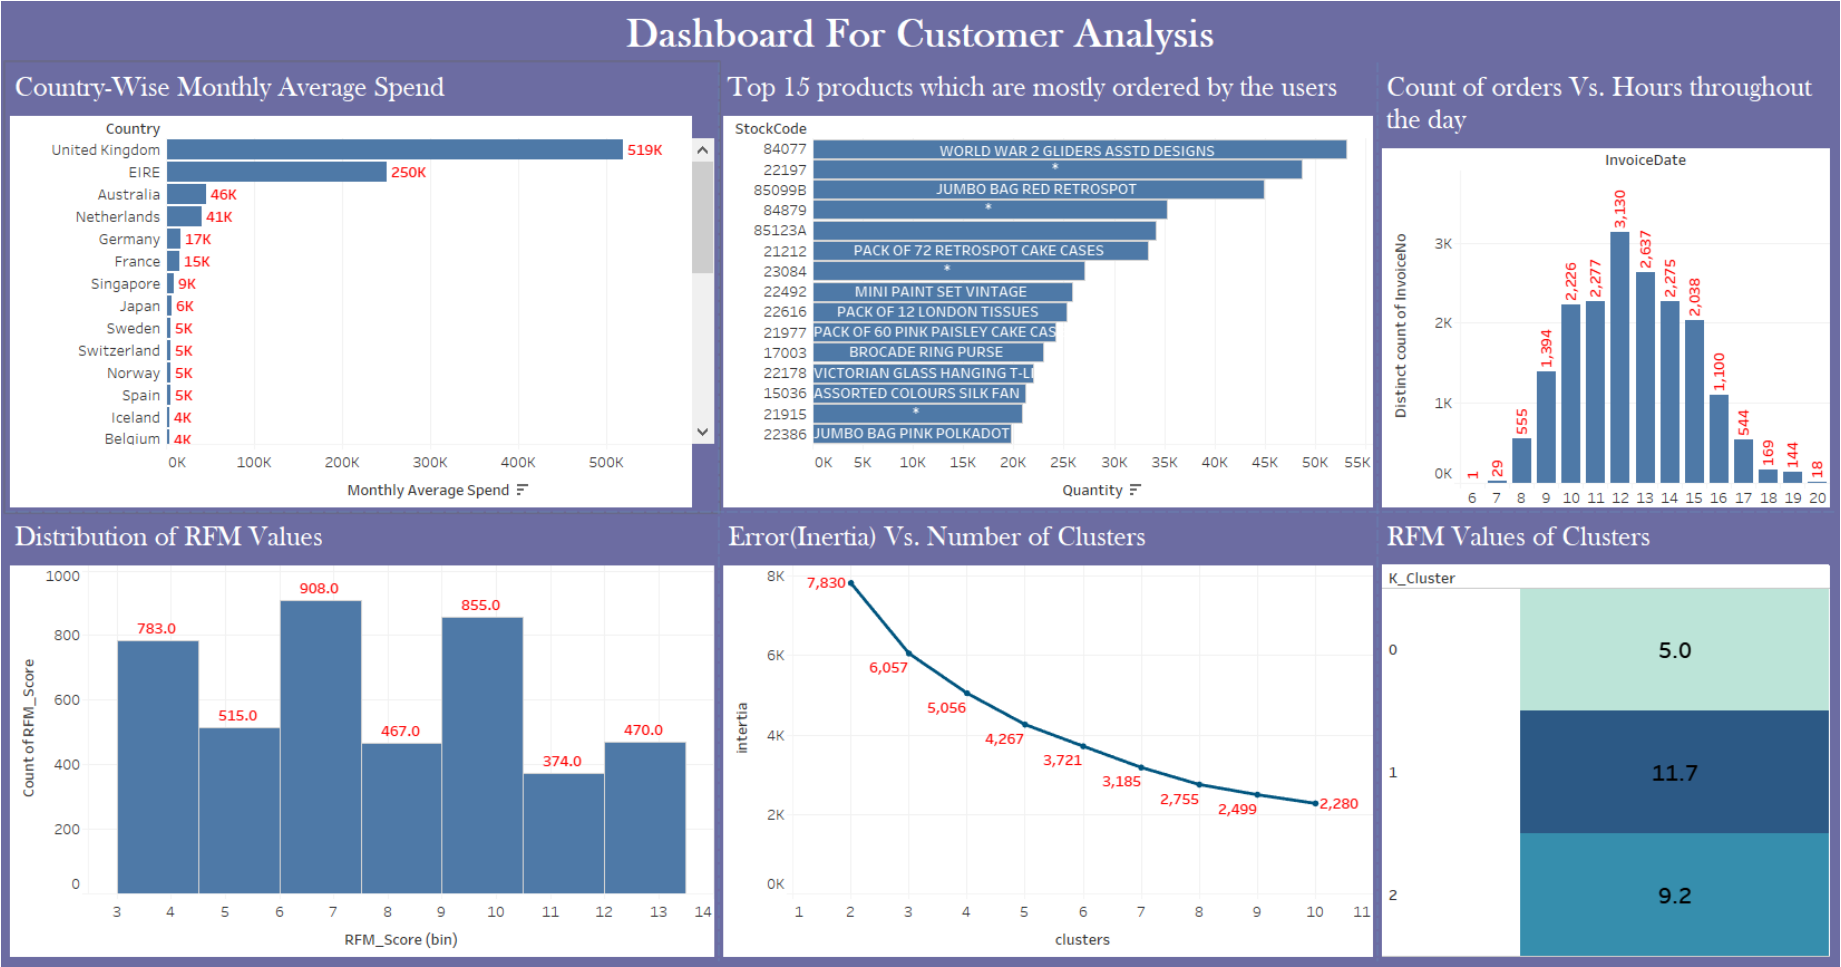
Error(Inertia) Vs. Number of Clusters



RFM Values of Clusters



Screenshot of Dashboard



Thank You

